

A Safety Net for Everyday AI

FTEC5660 | Personal Reflection | News published September 20-28, 2026

I already believed that AI would become part of everyday life. I can imagine asking an assistant to order milk tea, arrange appointments and handle routine purchases, much as I would ask a personal secretary. Four recent news stories have strengthened this belief, but they have also made my career interests more specific. If I work at the intersection of finance and AI, I want to help build the safety net that allows people to delegate tasks without giving up control over their money and information.

On September 26, TechCrunch reported that Google was testing purchases from Flipkart through Gemini and AI Mode in India [1]. Although this was a limited test, it made the idea of an everyday AI assistant feel more concrete to me. The assistant would be involved in a transaction, rather than only suggesting what to buy. That raises practical questions: what exactly did the customer authorize, what happens if an order is duplicated, and who handles an incorrect purchase? A small shopping mistake might be manageable, but the same ambiguity could have serious consequences in a financial transaction.

Meta announced its Enterprise Platform on September 28, bringing its AI technology to businesses and developers [2]. This announcement made me think about how personal assistants might connect with business systems. An assistant handling a request could interact with a retailer, a payment provider and other services. I would want its permission to remain clear throughout that process. Permission to compare prices should not automatically become permission to spend money. This is one reason I am interested in protocols that communicate the scope and limits of a user's authorization across services.

The scale of investment also caught my attention. Reuters reported on September 24 that Anthropic had agreed to a seven-year, US\$11.6 billion cloud-services commitment with Akamai [3]. This does not prove that every AI service will succeed, but it shows how substantial the infrastructure commitments behind expansion have become. It made me consider what happens when a system moves beyond a small demonstration. A mistake that is easy to notice in one transaction could be harder to detect across thousands. Reliability, monitoring and the cost of checking results would therefore need to be part of my design from the beginning.

Finally, Bloomberg reported on September 20 that Microsoft AI chief Mustafa Suleyman argued that competition with China should not justify abandoning AI safeguards [4]. I agree with the broader principle that stronger capabilities should be accompanied by clear limits. In finance, I would translate that principle into specific controls: restricted data access, spending limits, explicit confirmation for high-risk actions, and records of what the user authorized and what the agent actually did. When a request exceeds its authority, the system should stop and ask. Since some financial actions cannot easily be reversed, checks before execution would be particularly important.

These stories have not changed my belief that AI will enter everyday life; they have changed what I want to contribute to that future. I am interested in making AI useful within boundaries that people can understand and trust. In future coursework and projects, I would like to go beyond getting a model to produce an answer. I want to test unclear instructions, duplicate requests and unauthorized actions, and learn how to combine technical controls with financial risk management. A protocol alone cannot guarantee safety, but well-defined permissions, monitoring and human intervention could form a practical safety net. Building that connection between capability and responsibility is a direction I would like to explore in my career.

References

- [1] TechCrunch. (September 26, 2026). *Google tests buying from Walmart-owned Flipkart through Gemini and AI Mode in India*. [Read article](#).
- [2] Meta Newsroom. (September 28, 2026). *Launching Meta Enterprise Platform*. [Read article](#).
- [3] Reuters, via Investing.com. (September 24, 2026). *Akamai, Anthropic sign \$11.6 billion cloud services deal*. [Read article](#).
- [4] Bloomberg News, via Bloomberg Law. (September 20, 2026). *Microsoft AI Chief Says China Isn't Excuse to Forego Regulation*. [Read article](#).